

DataLens Pro Version 6.0: Democratizing Enterprise Business Intelligence Through Automated Machine Learning and Natural Language Processing in a Single-Page Application

Sayantan Roy

Computer Science and Engineering

Government College of Engineering and Ceramic Technology (GCECT), Kolkata

Email: sayantanr32@gmail.com

GitHub Repository: <https://github.com/sayantanr/DataLens-pro-version-6.0>

Abstract

The exponential growth of enterprise data has necessitated the development of highly agile, intuitive, and robust Business Intelligence (BI) platforms. Traditional BI solutions often suffer from steep learning curves, expensive licensing, reliance on cloud computing infrastructure, and rigid data pipeline prerequisites. In this paper, we introduce DataLens Pro Version 6.0, a groundbreaking, entirely open-source, Python-based single-page web application engineered to replace conventional enterprise BI tools like Tableau and PowerBI. Built on the Streamlit framework, DataLens Pro 6.0 dynamically ingests heterogeneous data formats including CSV, Excel, JSON, XML, and raw SQL scripts, automatically inferring data types to generate over 50 interactive visualizations instantaneously. Furthermore, the software pioneers a novel paradigm in automated analytics by autonomously synthesizing over 100 relevant synthetic metrics, Key Performance Indicators (KPIs), and Objectives and Key Results (OKRs). The platform integrates a comprehensive Automated Machine Learning (AutoML) workbench capable of training and evaluating over 15 distinct regression models, performing K-Means clustering, Isolation Forest anomaly detection, and ARIMA time-series forecasting. A dual-engine Natural

Language Processing (NLP) system enables conversational querying of data strictly offline without telemetry. This paper comprehensively details the architectural design, algorithmic implementation, and feature sets of DataLens Pro 6.0, demonstrating its efficacy in democratizing advanced data science capabilities for non-technical users while preserving strict data privacy through local execution.

1 Introduction

In the contemporary digital economy, data is the primary driver of strategic decision-making. Organizations collect vast amounts of information spanning sales, customer interactions, telemetry, and logistics. However, extracting actionable insights from this raw data remains a significant bottleneck. Standard operating procedures dictate that data engineers format data, data scientists build models, and business analysts construct dashboards. This tripartite process is inherently slow and resource-intensive. Proprietary Business Intelligence (BI) platforms such as Tableau, Microsoft PowerBI, and Looker have attempted to streamline this workflow. However, these tools are often accompanied by substantial financial barriers, restrictive data row limits, and a steep learning curve requiring specialized training (e.g., DAX, VizQL). More critically, these platforms increasingly rely on cloud-based processing, raising severe data privacy and sovereignty concerns for organizations dealing with highly sensitive information such as healthcare records (HIPAA) or financial transactions. To address these systemic inefficiencies and privacy concerns, we propose DataLens Pro Version 6.0. DataLens Pro is a monolithic, single-page application built entirely in Python using the Streamlit library. The software fundamentally redefines the BI workflow by prioritizing automation, offline execution, and universal accessibility. Upon the simple drag-and-drop of a dataset, DataLens Pro executes an automated feature extraction pipeline that instantly yields over 50 sophisticated, interactive dashboards. This research paper provides an exhaustive examination of the structural and functional components of DataLens Pro 6.0. We will explore the data ingestion engine, the automated visualization generation logic, the novel automated KPI and OKR synthesis algorithms, the integrated machine learning pipeline, and the offline natural language parsing architecture.

2 Architectural Design and Implementation

DataLens Pro 6.0 is architected as a lightweight, robust single-page application (SPA). By leveraging Streamlit, the application bypasses traditional web development complexities (HTML/CSS/JS DOM manipulation) in favor of a purely Pythonic reactive state-machine.

2.1 Universal Data Ingestion Engine

A primary limitation of existing BI tools is rigid data ingestion protocols. DataLens Pro 6.0 addresses this by implementing a highly versatile, caching-enabled data loader. The application natively supports multiple file formats through dynamic file signature recognition and parsing logic:

- **CSV and Excel (.xls, .xlsx):** Parsed using highly optimized Pandas reading functions, bypassing encoding errors automatically.
- **JSON and XML:** The application natively parses nested hierarchical data structures, flattening them into relational DataFrames for immediate analysis.
- **In-Memory SQL Execution:** A unique feature of DataLens Pro 6.0 is its ability to ingest raw '.sql' script files. The software establishes an isolated, in-memory SQLite database connection, executes the script, queries the resulting master tables, and seamlessly imports the structural data into the Pandas ecosystem.

2.2 Intelligent File Merging and Type Detection

When a user uploads multiple files simultaneously, the ingestion engine utilizes the Pandas concat function to merge the disparate datasets into a single cohesive global dataframe. Following ingestion, a deterministic algorithm executes to classify every column. This algorithm differentiates columns into Dimensions (categorical string or low-cardinality numeric data), Measures (high-cardinality continuous numeric data), and Dates (temporal data). This auto-classification forms the foundational logic for the automated dashboard generation.

3 Automated Feature Extraction and Visualization

The core philosophy of DataLens Pro 6.0 is zero-configuration analytics. The application dynamically generates tabs corresponding to different analytical domains. Each tab utilizes the Plotly library to render highly interactive, WebGL-accelerated charts.

3.1 Executive KPIs and Statistical Summaries

The initial view presents the user with high-level KPI cards. The application iterates through the identified Measures, calculating the global sum and mean, and presents these figures using gradient-styled metric components. Concurrently, a statistical summary matrix is generated, detailing the count, mean, standard deviation, and quartiles of the data.

3.2 Distribution and Categorical Analytics

Understanding data spread is critical for identifying skewness and outliers. DataLens Pro automates this by generating nested Box Plots, Violin Plots, and Histograms for every numerical measure. For categorical Dimensions, the software autonomously generates Bar Charts and Pie Charts, highlighting proportional frequency distributions without requiring user pivoting.

3.3 Temporal Analysis and Time-Series

If temporal data is detected, the application initiates the Time-Series module. It aggregates continuous numerical measures over month-year periods, overlaying Sum and Average calculations on dual-axis interactive line charts.

3.4 Correlation and Dimensionality Analytics

To uncover hidden collinearity, DataLens Pro 6.0 generates Pearson correlation matrices mapped to interactive heatmaps. For datasets with sufficient dimensionality, a Principal Component Analysis (PCA) pipeline is automatically executed. The software scales the data, reduces it to two principal components, and plots the resulting variance explanation on an interactive scatter graph, enabling rapid identification of multi-dimensional data clustering.

4 Automated Business Intelligence: Metrics, KPIs, and OKRs

A groundbreaking addition to DataLens Pro Version 6.0 is the Automated Business Intelligence generation algorithm. Traditional dashboards require users to mathematically define what they wish to track. DataLens Pro inverts this paradigm. Upon data upload, the software executes a synthesis algorithm that generates over 100 business management structures.

- **Synthetic Metrics:** The engine isolates numeric columns and systematically generates definitions for Sums, Means, Maximums, Minimums, and Distinct Counts.
- **KPI Synthesis:** The software converts the raw metrics into formatted Key Performance Indicators. It explicitly states the underlying mathematical formula (e.g., *SUM(Sales)*) and renders the exact current value.
- **OKR Extrapolation:** Taking the analytics a step further, the software generates theoretical Objectives and Key Results. If a measure represents "Profit", the software dynamically authors an Objective ("Improve Profit performance") and a Key Result ("Increase total Profit by 10%"), instantly calculating the theoretical target threshold.

These generated structures are rendered in distinct tables within a dedicated tab and can be instantly exported as CSV files for integration into broader corporate management software.

5 Machine Learning and Predictive Analytics Pipeline

DataLens Pro 6.0 democratizes Data Science by embedding a comprehensive Automated Machine Learning (AutoML) workbench directly into the application, abstracting away the complex scikit-learn API syntax.

5.1 The AutoML Leaderboard

Users select a target variable and an arbitrary number of predictor features. The software then splits the data, applies necessary encoding, handles missing values, and initiates a parallel training sequence across 15 distinct regression models. These models include:

- Linear, Ridge, Lasso, and ElasticNet Regression.
- Tree-based ensembles: Random Forest, Extra Trees, Gradient Boosting, AdaBoost, and Decision Trees.
- Support Vector Regression (SVR) and K-Neighbors Regressor.
- Multi-layer Perceptron (MLP) Neural Networks and Gaussian Process Regression.

Following training, the application dynamically generates a Leaderboard, ranking the algorithms based on R^2 Score, Mean Absolute Error (MAE), and Root Mean Squared Error (RMSE).

The top-performing algorithm is algorithmically crowned as the winner, drastically reducing the time required for model selection.

5.2 Clustering and Anomaly Detection

Beyond supervised regression, the platform offers unsupervised learning utilities.

- **K-Means Clustering:** Users can specify the desired number of clusters (K) via a slider. The application scales the selected measures and plots a 3D interactive scatter plot colored by the K-Means cluster assignments.
- **Isolation Forest:** For fraud detection or data cleansing, the application runs an Isolation Forest algorithm to partition data into Inliers and Outliers, visualizing the contamination distribution via an interactive pie chart.

5.3 ARIMA Forecasting

For temporal datasets, DataLens Pro incorporates an Auto-Regressive Integrated Moving Average (ARIMA) model. The software automatically frequency-aligns the time series, handles missing dates via forward-filling, and forecasts the subsequent 30 days of activity, rendering the predictive trajectory on a Plotly line chart.

6 Natural Language Processing and Smart Querying

DataLens Pro 6.0 recognizes that visual dashboards cannot account for every ad-hoc analytical query. Therefore, it incorporates a Smart Query Workbench allowing natural language interaction with the dataset.

6.1 Offline Smart Parser

To maintain the strict local-execution privacy mandate, the software includes a custom-built, regex-driven Offline Smart Parser. Users can input queries such as "average sales by region" or "filter where profit $\geq$ 500". The parser algorithmically extracts the requested mathematical operation, identifies the target measure, isolates the grouping dimension, and dynamically executes the corresponding Pandas groupby operations. It instantly renders both the tabular result

and an accompanying Plotly bar chart, all without requiring an internet connection or external API calls.

6.2 PandasAI Integration

For users who require highly complex, conversational queries and are willing to utilize cloud APIs, the software optionally integrates PandasAI. By providing an OpenAI API key, users can leverage Large Language Models (LLMs) to perform intricate data manipulations, aggregations, and inferential analyses simply by asking questions in plain English.

7 Advanced Statistical Validation

DataLens Pro 6.0 caters to rigorous academic and statistical requirements by including a dedicated Statistical Insights tab.

- **Quantile-Quantile (QQ) Plots:** The application utilizes Scipy to automatically generate QQ plots for numerical measures, allowing statisticians to visually verify normality assumptions against theoretical normal distributions.
- **Automated ANOVA Testing:** The software automatically pairs categorical dimensions with numerical measures and utilizes statsmodels to run Analysis of Variance (ANOVA) tests. It instantly renders the resulting F-statistics and P-values, enabling rapid identification of statistically significant variances across categorical groups.

8 Security, Privacy, and Global Filtering Architecture

8.1 Data Sovereignty

A critical advantage of DataLens Pro 6.0 over its commercial counterparts is its execution environment. Because the application runs entirely via a local Python process (`streamlit run z.py`), data never leaves the host machine. There is zero telemetry, tracking, or cloud uploading required, ensuring total compliance with stringent corporate and governmental data privacy regulations.

8.2 Global Filtering and State Management

The application implements a global filtering architecture housed within the sidebar. When a user adjusts a slider or multiselect box, the global Pandas dataframe is filtered. Because of Streamlit's reactive rendering paradigm, this single state change cascades downwards, instantly updating all 50+ dashboards, statistical tests, and machine learning models to reflect the newly filtered dataset.

8.3 Enterprise Exportation

All analysis generated within DataLens Pro 6.0 is highly portable. Users can export the filtered dataframes to CSV. Furthermore, the application can compile the data previews, KPI tables, and critical plots into a standalone, statically generated HTML report, allowing insights to be shared easily across an organization without requiring the recipient to have the software installed.

9 Future Work and Enhancements

While DataLens Pro Version 6.0 represents a massive leap forward in open-source BI, several avenues for future research and enhancement exist. Subsequent iterations aim to incorporate deeper Deep Learning frameworks (such as PyTorch and TensorFlow) for more complex non-linear time series forecasting (e.g., LSTMs). Additionally, the implementation of a distributed computing backend (such as Dask or PySpark) would allow the application to effortlessly scale to multi-terabyte datasets that currently exceed standard local memory limits.

10 Conclusion

DataLens Pro Version 6.0 fundamentally disrupts the traditional Business Intelligence landscape. By amalgamating a Universal Data Ingestion Engine, fully automated dashboard generation, autonomous KPI and OKR synthesis, and a complete Machine Learning workbench into a single, locally hosted Python application, it eliminates the historical barriers of cost, complexity, and privacy concerns. The software proves that advanced data analytics and predictive modeling can be fully democratized and made accessible to anyone through intuitive design and robust automation algorithms. As open-source frameworks continue to evolve, platforms like

DataLens Pro will undoubtedly become the standard for agile, secure, and intelligent enterprise data analysis.